

# assignment0

October 7, 2023

## 1 CS260R Reinforcement Learning Assignment 0: Jupyter Notebook usage and assignment submission workflow

---

*CS260R 2023Fall: Reinforcement Learning. Department of Computer Science at University of California, Los Angeles. Course Instructor: Professor Bolei ZHOU. Assignment author: Zhenghao PENG, Yiran WANG.*

You are asked to finish four tasks:

1. Fill in your name and University ID in the next cell.
2. Install pytorch and finish the Kindergarten Pytorch section.
3. Run all cells and save this notebook as a **PDF** file.
4. Compress this folder **assignment0** as a **ZIP** file and submit the **PDF** file and the **ZIP** file separately as two files in BruinLearn.

```
[1]: # TODO: Fill your name and UID here
my_name = "Haniyeh Ehsani Oskouie"
my_student_id = "306300374"
```

```
[2]: # Run this cell without modification

text = "Oh, I finished this assignment! I am {} ({}).".format(my_name,
↳my_student_id)
print(text)
with open("{} .txt".format(text), "w") as f:
    f.write(text)
```

Oh, I finished this assignment! I am Haniyeh Ehsani Oskouie (306300374)

### 1.1 Kindergarten Pytorch

1. Please install pytorch in your virtual environment following the instruction: <https://pytorch.org/get-started/locally/>.

pip install torch torchvision

2. If you are not familiar with Pytorch, please go through [the tutorial in official website](#) until you can understand the [quick start tutorial](#).

3. The following code is copied from the [quick start tutorial](#), please solve all TODOs and print the result in the cells before generating the PDF file.

### 1.1.1 Prepare data

```
[3]: import torch
from torch import nn
from torch.utils.data import DataLoader
from torchvision import datasets
from torchvision.transforms import ToTensor

# Download training data from open datasets.
training_data = datasets.FashionMNIST(
    root="data",
    train=True,
    download=True,
    transform=ToTensor(),
)

# Download test data from open datasets.
test_data = datasets.FashionMNIST(
    root="data",
    train=False,
    download=True,
    transform=ToTensor(),
)

batch_size = 64

# Create data loaders.
train_dataloader = DataLoader(training_data, batch_size=batch_size)
test_dataloader = DataLoader(test_data, batch_size=batch_size)
```

Downloading <http://fashion-mnist.s3-website.eu-central-1.amazonaws.com/train-images-idx3-ubyte.gz>

Downloading <http://fashion-mnist.s3-website.eu-central-1.amazonaws.com/train-images-idx3-ubyte.gz> to data/FashionMNIST/raw/train-images-idx3-ubyte.gz

100.0%

Extracting data/FashionMNIST/raw/train-images-idx3-ubyte.gz to data/FashionMNIST/raw

Downloading <http://fashion-mnist.s3-website.eu-central-1.amazonaws.com/train-labels-idx1-ubyte.gz>

Downloading <http://fashion-mnist.s3-website.eu-central-1.amazonaws.com/train-labels-idx1-ubyte.gz> to data/FashionMNIST/raw/train-labels-idx1-ubyte.gz

100.0%

Extracting data/FashionMNIST/raw/train-labels-idx1-ubyte.gz to  
data/FashionMNIST/raw

Downloading http://fashion-mnist.s3-website.eu-  
central-1.amazonaws.com/t10k-images-idx3-ubyte.gz  
Downloading http://fashion-mnist.s3-website.eu-  
central-1.amazonaws.com/t10k-images-idx3-ubyte.gz to  
data/FashionMNIST/raw/t10k-images-idx3-ubyte.gz

100.0%

Extracting data/FashionMNIST/raw/t10k-images-idx3-ubyte.gz to  
data/FashionMNIST/raw

Downloading http://fashion-mnist.s3-website.eu-  
central-1.amazonaws.com/t10k-labels-idx1-ubyte.gz  
Downloading http://fashion-mnist.s3-website.eu-  
central-1.amazonaws.com/t10k-labels-idx1-ubyte.gz to  
data/FashionMNIST/raw/t10k-labels-idx1-ubyte.gz

100.0%

Extracting data/FashionMNIST/raw/t10k-labels-idx1-ubyte.gz to  
data/FashionMNIST/raw

### 1.1.2 Define model

```
[6]: # Get cpu, gpu or mps device for training.
device = (
    "cuda"
    if torch.cuda.is_available()
    else "mps"
    if torch.backends.mps.is_available()
    else "cpu"
)
print(f"Using {device} device")

# Define model
class NeuralNetwork(nn.Module):
    def __init__(self):
        super().__init__()
        self.flatten = nn.Flatten()

        # TODO: Define the self.linear_relu_stack by uncommenting next few lines
        # and understand what they mean
        self.linear_relu_stack = nn.Sequential(
            nn.Linear(28*28, 512),
```

```

        nn.ReLU(),
        nn.Linear(512, 512),
        nn.ReLU(),
        nn.Linear(512, 10)
    )

    def forward(self, x):
        x = self.flatten(x)
        logits = self.linear_relu_stack(x)
        return logits

model = NeuralNetwork().to(device)
print(model)

```

Using mps device

```

NeuralNetwork(
  (flatten): Flatten(start_dim=1, end_dim=-1)
  (linear_relu_stack): Sequential(
    (0): Linear(in_features=784, out_features=512, bias=True)
    (1): ReLU()
    (2): Linear(in_features=512, out_features=512, bias=True)
    (3): ReLU()
    (4): Linear(in_features=512, out_features=10, bias=True)
  )
)

```

### 1.1.3 Define training and test pipelines

```

[7]: loss_fn = nn.CrossEntropyLoss()
optimizer = torch.optim.SGD(model.parameters(), lr=1e-3)

def train(dataloader, model, loss_fn, optimizer):
    size = len(dataloader.dataset)
    model.train()
    for batch, (X, y) in enumerate(dataloader):
        X, y = X.to(device), y.to(device)

        # Compute prediction error
        pred = model(X)
        loss = loss_fn(pred, y)

        # Backpropagation

        # TODO: Uncomment next three lines and understand what they mean
        loss.backward()
        optimizer.step()
        optimizer.zero_grad()

```

```

        if batch % 100 == 0:
            loss, current = loss.item(), (batch + 1) * len(X)
            print(f"loss: {loss:>7f} [{current:>5d}/{size:>5d}]")

def test(dataloader, model, loss_fn):
    size = len(dataloader.dataset)
    num_batches = len(dataloader)
    model.eval()
    test_loss, correct = 0, 0
    with torch.no_grad():
        for X, y in dataloader:
            X, y = X.to(device), y.to(device)
            pred = model(X)
            test_loss += loss_fn(pred, y).item()

            # TODO: Uncomment next line and understand what it means
            correct += (pred.argmax(1) == y).type(torch.float).sum().item()

    test_loss /= num_batches
    correct /= size
    print(f"Test Error: \n Accuracy: {(100*correct):>0.1f}%, Avg loss: \n
    ↳{test_loss:>8f} \n")

```

#### 1.1.4 Run the training and test pipelines

```

[8]: epochs = 5
    for t in range(epochs):
        print(f"Epoch {t+1}\n-----")
        train(train_dataloader, model, loss_fn, optimizer)
        test(test_dataloader, model, loss_fn)
    print("Done!")

```

Epoch 1

```

-----
loss: 2.303508 [ 64/60000]
loss: 2.290811 [ 6464/60000]
loss: 2.275872 [12864/60000]
loss: 2.261932 [19264/60000]
loss: 2.243488 [25664/60000]
loss: 2.210532 [32064/60000]
loss: 2.223848 [38464/60000]
loss: 2.187051 [44864/60000]
loss: 2.183177 [51264/60000]
loss: 2.143207 [57664/60000]
Test Error:
Accuracy: 40.6%, Avg loss: 2.144574

```

Epoch 2

```
-----  
loss: 2.155992 [ 64/60000]  
loss: 2.146238 [ 6464/60000]  
loss: 2.091598 [12864/60000]  
loss: 2.100174 [19264/60000]  
loss: 2.045278 [25664/60000]  
loss: 1.976297 [32064/60000]  
loss: 2.006233 [38464/60000]  
loss: 1.925260 [44864/60000]  
loss: 1.928890 [51264/60000]  
loss: 1.844805 [57664/60000]
```

Test Error:

Accuracy: 55.8%, Avg loss: 1.855037

Epoch 3

```
-----  
loss: 1.881813 [ 64/60000]  
loss: 1.852645 [ 6464/60000]  
loss: 1.745204 [12864/60000]  
loss: 1.784674 [19264/60000]  
loss: 1.669344 [25664/60000]  
loss: 1.616550 [32064/60000]  
loss: 1.640001 [38464/60000]  
loss: 1.548512 [44864/60000]  
loss: 1.576994 [51264/60000]  
loss: 1.463080 [57664/60000]
```

Test Error:

Accuracy: 62.4%, Avg loss: 1.494204

Epoch 4

```
-----  
loss: 1.550790 [ 64/60000]  
loss: 1.519951 [ 6464/60000]  
loss: 1.384457 [12864/60000]  
loss: 1.455961 [19264/60000]  
loss: 1.343319 [25664/60000]  
loss: 1.329989 [32064/60000]  
loss: 1.345554 [38464/60000]  
loss: 1.276705 [44864/60000]  
loss: 1.318559 [51264/60000]  
loss: 1.207958 [57664/60000]
```

Test Error:

Accuracy: 63.6%, Avg loss: 1.243244

Epoch 5

```
-----  
loss: 1.310662 [ 64/60000]
```

```
loss: 1.295911 [ 6464/60000]
loss: 1.141164 [12864/60000]
loss: 1.244906 [19264/60000]
loss: 1.131477 [25664/60000]
loss: 1.141375 [32064/60000]
loss: 1.163275 [38464/60000]
loss: 1.104242 [44864/60000]
loss: 1.154141 [51264/60000]
loss: 1.056673 [57664/60000]
```

Test Error:

Accuracy: 64.8%, Avg loss: 1.084810

Done!

### 1.1.5 Save model

```
[9]: torch.save(model.state_dict(), "model.pth")
print("Saved PyTorch Model State to model.pth")
```

Saved PyTorch Model State to model.pth

### 1.1.6 Load model and run the inference

```
[10]: model = NeuralNetwork().to(device)
model.load_state_dict(torch.load("model.pth"))
```

[10]: <All keys matched successfully>

```
[11]: classes = [
    "T-shirt/top",
    "Trouser",
    "Pullover",
    "Dress",
    "Coat",
    "Sandal",
    "Shirt",
    "Sneaker",
    "Bag",
    "Ankle boot",
]

model.eval()
x, y = test_data[0][0], test_data[0][1]
with torch.no_grad():
    x = x.to(device)
    pred = model(x)
    predicted, actual = classes[pred[0].argmax(0)], classes[y]
    print(f'Predicted: "{predicted}", Actual: "{actual}"')
```

Predicted: "Ankle boot", Actual: "Ankle boot"